FREE ARTICLE Explore Premium Services

EARNINGS CALL TRANSCRIPT

IBM (IBM) Q1 2026 Earnings Call Transcript

By Motley Fool Transcribing – Updated Apr 22, 2026 at 7:45PM EST

Fool.com on G Share Summarize with AI

Advertisement

Image source: The Motley Fool.

Advertisement
STOCKS MENTIONED
International Business Machines NYSE: IBM $290.49 (+1.03%) +$2.93
Motley Fool Stock Advisor's Latest Pick Get Access
*Average returns of all recommendations since inception. Cost basis and return based on previous market day close.
Advertisement
Advertisement

Date

Wednesday, April 22, 2026, at 5 p.m. ET

Call participants

- Chief Executive Officer — Arvind Krishna
- Senior Vice President and Chief Financial Officer — James J. Kavanaugh
- Vice President of Investor Relations — Olympia McNerney

Need a quote from a Motley Fool analyst? Email pr@fool.com

Takeaways

- **Revenue growth** -- Grew 6%, reflecting balanced performance across segments and the durability of the company's portfolio.
- **Operating pretax margin** -- Expanded 140 basis points, indicating higher profitability and efficiency.
- **Free cash flow** -- Rose 13% to $2.2 billion, representing the highest first-quarter free cash flow in a decade and a record margin.
- **Software revenue** -- Increased 8%, with annual recurring revenue at $24.6 billion, up 10%.
- **Data revenue** -- Climbed 16%, powered by generative AI demand and incremental contribution from Confluent and data stack acquisitions.
- **Red Hat** -- Delivered 10% growth, with OpenShift surpassing $2 billion in ARR, and virtualization contracts exceeding $600 million since early 2024.
- **Infrastructure revenue** -- Rose 12%; IBM Z posted 48% growth, with distributed infrastructure increasing at a double-digit rate.
- **Consulting revenue** -- Grew 1%, with signings rising 6%, and generative AI comprising about 30% of the backlog.
- **Productivity savings** -- Achieved $4.5 billion since 2023, expecting an additional $1 billion in 2026, driven by AI-enabled transformation.
- **Segment margins** -- Infrastructure margin expanded 720 basis points, software margin up 60 basis points, consulting margin declined due to currency headwinds and reinvestment.
- **Strategic M&A activity** -- Invested $10.5 billion in acquisitions with the closing of Confluent, contributing over 15 points to the projected 20%-25% full-year data revenue growth.
- **AI integration** -- AI penetrates all business segments, including software, infrastructure, and consulting, with software AI revenue exceeding $1.5 billion, growing over 40% annually.
- **Mainframe AI adoption** -- IBM systems perform 450 billion inferences daily at 1 millisecond response time, supporting real-time fraud detection and new monetization streams.
- **Quantum initiatives** -- Remain on track to deliver a large-scale fault-tolerant quantum computer by 2029; recent demonstrations include pharmaceutical and material simulations.
- **Guidance** -- Management reaffirmed constant currency revenue growth of 5%+ and $1 billion free cash flow growth for the full year, with software growth outlook boosted to 10%+.

Summary

IBM (IBM +1.02%) articulated its continued strategic shift toward software-led hybrid cloud and AI platforms, substantiated by margin expansion and double-digit free cash flow growth. Executives highlighted significant integration of generative AI across all segments, emphasizing data and Red Hat as major growth drivers. The closing of Confluent is expected to be highly accretive, materially boosting data revenue by more than 15 percentage points within the projected 20%-25% range for the year. Operational discipline and AI-driven transformation led to sustained productivity gains and improved segment margins, except in consulting, which faced minor currency and reinvestment impacts. Mainframe innovation, notably in AI inferencing and modernization, underpins new monetization opportunities, including high-volume transaction processing and associated software uptake. Quantum computing progress was evidenced by milestone collaborations and plans to deliver a fault-tolerant system by 2029.

- Management described IBM's commercial focus as "enabling software," positioning the portfolio to capitalize on structural AI demand tailwinds across infrastructure, automation, and data layers.
- Software's annuity base reached $24.6 billion in annual recurring revenue, with high-value, recurring contracts now comprising 80% of the $30 billion trailing-twelve-month software line.
- AI-related software and consulting revenues are accelerating, with AI platform, agent, and orchestration offerings now representing over $1.5 billion and about 25% penetration of IBM's software business, while generative AI drives approximately 30% of consulting backlog.
- IBM Z's hardware placement grew by more than $1 billion in its first year, supporting a projected 3x-4x software and services stack revenue multiplier for future cycles.
- IBM Bob, the internal AI-driven software development platform, is now GA and shows 45% average developer productivity improvements, contributing to operating efficiency.
- Management committed to maintaining a "neutral and Switzerland-like" approach to AI model integration, enabling clients to deploy any model — open-source or proprietary — on any infrastructure, regardless of ultimate model winner.
- Leadership cited strong, balanced geographic performance; macro uncertainty exists, but has not yet affected deal progression or pipelines, except for the potential of future supply chain or energy shocks.
- Consulting signings momentum is driven by large GenAI transformation deals, with 80% of GenAI consulting business now tied to new clients, and ARR surpassing $4 billion for the segment.

Industry glossary

- **ARR (Annual Recurring Revenue)**: The predictable, contracted revenue IBM expects to collect annually from subscription- or consumption-based agreements.
- **IBM Z**: IBM's enterprise mainframe platform, designed for ultra-high availability, security, and transaction volume, now including AI inferencing capabilities.
- **MIPS (Million Instructions Per Second)**: A capacity metric used for computing power on IBM mainframes, tied to billing and performance measurement.
- **GenAI (Generative AI)**: Artificial intelligence models and solutions capable of generating content, predictions, or automating processes, deployed across IBM's products and services.
- **OpenShift**: A Red Hat-developed, Kubernetes-based container platform for running applications consistently across hybrid cloud and on-premise environments.
- **watsonx**: IBM's platform for AI development and orchestration, integrating tools for data, automation, security, and multi-model workflows.
- **Confluent**: An acquisition providing IBM with real-time streaming data infrastructure, facilitating governed data feeds for AI and analytics.
- **ARM**: A semiconductor and software design company; in this context, refers to IBM's strategic initiative to integrate ARM-based software ecosystems into its infrastructure offerings.
- **Quantum Advantage**: The milestone when a quantum computer demonstrably outperforms classical computing techniques in practical applications.

Full Conference Call Transcript

Arvind Krishna: Thank you for joining us today. Let me start with our first quarter results and then provide context on what we are seeing across the business. IBM is off to a strong start to 2026. Revenue in the first quarter grew 6% and combined with strong margin expansion, drove 13% growth in free cash flow. These results reflect the durability of our portfolio, the mission-critical nature of the work we do for clients and the continued execution of our strategy. Let me first touch on the macro. While we are operating in a dynamic environment, Middle East developments didn't impact us in the first quarter.

Uncertainties remain, but our diversity across businesses, geographies, industries and large enterprise clients position us well. Conversations we are having with clients remain consistent. Enterprises are investing in capabilities that increase resiliency, productivity and accelerate growth. They are modernizing core systems. They are scaling AI and they're making deliberate choices about where workloads should run and who controls the infrastructure underneath them. These are structural priorities and they align directly with IBM's strengths. This quarter's performance reinforces the strategic choices we have made over the last several years to advance IBM as a software-led hybrid cloud and AI platform company. Software revenue grew 8%, with Data and Red Hat growing double digits.

Infrastructure grew 12% with another record Z quarter, up 48%. We also had strong performance in Distributed Infrastructure as generative AI increases demand for our storage offerings. Consulting grew 1% with momentum in enterprise data and business application transformations as clients modernize to deploy AI securely and at scale. The durability of our portfolio is a defining feature of IBM today. Let me spend a few minutes on AI. Enterprises are still figuring out where to deploy this technology and where competitive advantage truly sits. Every major technology wave has followed a pattern. Value begins with infrastructure, moves to enabling platforms and ultimately concentrates in the workflows where businesses operate. Right now, the spotlight is on foundation models.

Enterprises are building portfolios, frontier models for some workloads, smaller models running on-premise for others and open source models where control and flexibility matter the most. Enterprises will want to retain control of the proprietary data. AI will run everywhere across public cloud, private and sovereign clouds and on-premise. The core challenge is making all of this work together. This includes orchestrating across models, agents and workflows, governing enterprise data and securing these systems at scale. And that is exactly where IBM operates. We are building the platform that lets enterprises put AI to work on their terms, wherever it runs, whichever models they choose and under governance they control.

Our portfolio is built around world-class security, support and integration for an enterprise environment. Red Hat provides a common open platform that lets enterprises run applications in AI consistently across any infrastructure. More AI adoption means more demand for open flexible infrastructure. In automation, the logic is similar, agents multiply applications, integrations and execution paths. Managing that sprawl requires a controlled plane to provision infrastructure, integrate applications, secure environments and manage cost. This is what our end-to-end automation portfolio provides. In an AI-driven world, security risks are rising. IBM Concert identifies vulnerabilities proactively and automates remediation, helping enterprises maintain resilience at scale. Our data business is seeing similar AI tailwinds.

AI is only as good as the data it can access. And increasingly, that data is not static. It is generated continuously across transactions, applications and interactions. To deliver real-term AI outcomes, data must be available in motion, governed and delivered securely to models and agents wherever they are running. Confluent, which we closed this past quarter, solves that directly. It streams live, governed data to models and agents across the hybrid environment. And the orchestration layer ties it together. In a multi-model world, clients need to route between models, manage agent workflows and maintain governance. That is what watsonx Orchestrate and our watsonx platform deliver.

We have also created AI additions of critical software products like Db2, Cognos and MQ. These embed agentic AI that can reason, act and automate at scale while preserving IBM enterprise-grade security and trust. Infrastructure remains a critical differentiator as AI moves into the core of enterprise operations. IBM Z delivers the lowest unit cost architecture at scale for workloads that require end-to-end encryption, continuous availability and ultra-high throughput. Clients rely on our Z platform to process billions of transactions reliably with 6 to 8 9s of availability. They run AI inferencing directly in line with those transactions.

Our Spyre accelerator lets clients run AI on 100% of the transaction volume without moving data off platform, allowing them to embed AI directly into their transaction flows. Financial services clients are using this for real-time fraud detection, saving tens of millions of dollars. At the same time, AI-assisted modernization, including code understanding, refactoring and API integration makes it easier to evolve applications without compromising the guarantees the platform provides. Our watsonx Assistant for Z made available over 2 years ago to help clients preserve the architectural strengths that deliver resilience, security and scalability while making the platform more productive. Clients who have deployed watsonx Code Assistant for Z are growing MIPS capacity 3x faster than those who have not.

In consulting, AI is both a growth driver and a productivity engine. As agents take on more work, delivery becomes faster, more software driven and more scalable. IBM is leading into the shift through our consulting advantage platform and unique integrated value sitting side-by-side with software. This helps clients operationalize AI while improving our own efficiency. Demand continues to accelerate as clients move beyond experimentation and focus on transforming applications, data and workflows to embed AI into core operation. All of this allows us to drive to drive value for clients. ServiceNow is leveraging watsonx for automated data quality and observability to deliver AI-ready data and code generation to refresh legacy applications to modern application run times, including ServiceNow.

Visa continues to work with IBM on ongoing software and data modernization efforts supporting the scale, resiliency and performance of VisaNet. With Nestle, we are using NVIDIA accelerated watsonx.data to embed AI directly into core order-to-cash operations, enabling faster real-time insights across Nestle's global supply chain. This highlights how quickly we can bring research to bear for commercial value. Nestle was ideal for this proof of concept because of a strong digital backbone. In infrastructure, clients such as NatWest and RBC are modernizing their mainframe environments using AI and automation capabilities, including watsonx Assistant and watsonx Code Assistant for Z to improve resiliency, security and developer productivity. We continue to accelerate organic innovation.

IBM Bob, our AI-based software development system, is now generally available. Our entire developer workforce is using Bob with average productivity gains of 45%. Bob automates the full software life cycle from legacy modernization to security using specialized agents and multimodal optimization. It drives developer productivity and predictable costs. We also introduced Sovereign Core software that lets organizations run AI workloads under their own operational authority within a defined jurisdiction and auditable controls. [PC] sovereignty becoming a bigger factor in where and how workloads run. Every enterprise and every nation is waking up to the same reality. They need AI and cloud infrastructure they control, infrastructure no one can turn off or tamper with because of geopolitics.

During the quarter, we also announced strategic collaborations with NVIDIA, expanding our work across GPU native analytics. In addition, we announced a strategic collaboration with ARM to expand how AI workloads run across IBM infrastructure. By enabling the ARM software ecosystem within mission-critical environments like IBM Z, clients can scale AI closer to the data while preserving the security and resilience they require. These partnerships reflect our approach, open and flexible and all the infrastructure clients choose. We continue to make progress in Quantum and remain on track to deliver the first, large-scale fault-tolerant quantum computer by 2029. Here are some signposts of progress.

In March, researchers used IBM Quantum hardware to simulate a 300 atom system with the Cleveland Clinic, demonstrating that quantum computers can serve as reliable tools for pharmaceutical discovery. Another team accurately simulated real magnetic materials. Magnetism is central to new forms of energy and electrification. These are significant demonstrations to date that quantum computers can serve as reliable tools for scientific discovery. We also released a new blueprint for Quantum-centric supercomputing that outlines the architecture for integrating quantum and classical systems at scale. We strongly believe that our partners will achieve the first examples of Quantum Advantage this year leveraging IBM hardware. In closing, we are executing on our strategy of accelerating revenue growth and delivering higher profitability.

Given our strong start to the year, we remain confident in our ability to sustain revenue growth of 5% plus and grow free cash flow by about $1 billion this year. With that, let me hand it over to Jim to go through the financials.

James Kavanaugh: Thanks, Arvind. In the first quarter, we delivered 6% revenue growth, 140 basis points of operating pretax margin expansion, 17% adjusted EBITDA growth, 19% diluted operating earnings per share growth and $2.2 billion of free cash flow, growing 13% year-to-year, representing our highest first quarter free cash flow in a decade and free cash flow margin in reported history. This performance reflects the work we have done to strengthen our software-led platforms, deliver innovation, value to clients and the durability of our financial model. Now I'll dive deeper into our segment performance. Software revenue grew 8% marking a strong start to the year.

This reflects the diversity of our portfolio, ongoing GenAI innovation, continued shift to higher-growth end markets and flexible consumption model. Our ARR was solid at $24.6 billion, up 10% since last year. Data revenue grew 16%, fueled by demand for our GenAI products, strength our strategic partnerships and inorganic contribution from data stack and Confluent, which closed in mid-March. Red Hat growth accelerated 2 points sequentially to 10%, largely driven by the stabilization of consumption-based services revenue growth that we expected. OpenShift is now $2 billion ARR business with strong growth. And virtualization continues to gain traction with over $600 million of contracts signed since the beginning of 2024.

Automation grew 7%, with February marking the 1-year anniversary of the acquisition of HashiCorp. Over the last year, we have seen record HashiCorp bookings, leveraging IBM's go-to-market scale and achieved adjusted EBITDA accretion ahead of expectations. Transaction processing grew again, up 2% as we monetize on a strong Z17 program. In infrastructure, our revenue grew 12% this quarter, with hybrid infrastructure up 25% and infrastructure support down 6%. Within hybrid infrastructure, growth was broad-based with strong demand for our offerings across IBM Z, Power and Storage. IBM Z continues to outperform prior programs, growing 48% this quarter.

Clients are investing in IBM Z as they modernize mission-critical workloads driven by requirements for resiliency, security and compliance, while enabling new AI capabilities on the platform. Distributed infrastructure grew double digits with strength in both power and storage. Our growth was driven by demand for [Power11] with its resiliency and performance advantages supporting data-intensive workloads. In storage, growth reflected strong adoption of our new flash offerings introduced in the first quarter, which incorporate industry-leading agentic AI capabilities. In consulting, our revenue grew 1% this quarter, reflecting momentum in the business as client demand continues to shift towards enterprise-wide transformation.

Signings returned to growth, up 6% with strength across our application and data transformation offerings, driven by clients modernizing their environments to support AI adoption and capture value. Revenue growth was balanced across the portfolio with both strategy and technology and intelligent operations up 1%. Generative AI is now firmly integrated across our consulting engagements, representing about 30% of our backlog. This reflects how generative AI has become embedded in the work we do. Our differentiated asset-led delivery model continues to drive productivity and speed to value, combining deep domain expertise with software automation and reusable assets to help clients deploy AI securely and at scale. Let me now discuss profitability.

Several years ago, we set an ambitious objective to reinvent our enterprise operations for greater speed, lower friction and structurally lower cost. Through disciplined execution, eliminating manual touch points, simplifying processes and applying data, automation and AI at scale, we have built a proven repeatable AI-enabled transformation engine that is accelerating. Since 2023, this has driven $4.5 billion of productivity savings, with an additional $1 billion expected in 2026. Our success is enabling us to accelerate investments in innovation, strengthen our competitive advantage as [client zero] and fuel our growth flywheel while expanding our margins.

You can see this in the results this quarter with productivity revenue scale and mix driving expansion of operating gross profit margin by 110 basis points, adjusted EBITDA margin by 170 basis points and operating pretax margin by 140 basis points, all ahead of expectations. Segment profit margins expanded by 720 basis points in infrastructure and 60 basis points in software. Consulting segment profit margin declined modestly, reflecting currency headwinds from geographic mix of the business and the reinvestment of productivity gains amid an improving demand environment. In the quarter, we generated $2.2 billion of free cash flow, up about $300 million year-over-year.

The primary driver of this growth is adjusted EBITDA, up about $600 million year-over-year, partially offset by higher net interest expense and increased investments in CapEx as we expected coming into 2026. We exited the first quarter with a strong liquidity position and a solid investment-grade balance sheet with cash of $11.8 billion. We invested $10.5 billion in acquisitions, driven by the closing of Confluent and returned $1.6 billion to shareholders in the form of dividends. Our debt balance ending the quarter was $66.4 billion, including debt of $12.8 billion for our financing business, with the receivables portfolio that is 80% in investment grade. Let me now pivot to discuss our expectations going forward.

The strong start to the year drives our confidence in delivering constant currency revenue growth of 5-plus percent in 2026 and free cash flow growth of about $1 billion year-over-year. Given where we are in the year, we believe it is prudent to maintain our guidance even as the underlying performance and execution are off to an encouraging start. The combination of our focused portfolio, investment in innovation and our diversity across businesses drives the durability of our performance. Our revenue expectations are underpinned by our accelerating software business, which we now expect to grow 10-plus percent this year.

In consulting, the quality of our backlog and momentum in GenAI with backlog penetration at about 30%, continue to support an acceleration in revenue growth to low to mid-single digits for the year. We are off to a great start with z17. And 4 quarters into z17's launch, we prudently continue to expect infrastructure revenue to be down low single digits for the year, representing about a 0.5 point impact to IBM. We remain confident this will be our strongest cycle given the AI innovation value we are delivering to clients. The momentum in our productivity flywheel is fueling margin expansion, while enabling investment in innovation.

Last quarter, we disclosed that we anticipated absorbing about $600 million of dilution from Confluent in 2026, driven largely by stock-based compensation and interest expense. While we are absorbing incremental dilution given the early closing of Confluent, actions we are taking to accelerate our cost synergies enable us to stay on track to expand operating pretax margins by about 1 point this year. Our operating tax rate for the year should be in the mid-teens and the timing of discrete items can cause the rate to vary within the year. For free cash flow, we continue to expect to grow about $1 billion for the full year, driven primarily by growth in adjusted EBITDA.

The headwinds I discussed heading into the year of higher cash taxes, higher CapEx and higher net interest expense remain the same. Looking to the second quarter, we expect our constant currency revenue growth rate to be similar to the full year. And for operating pretax margin, we expect about 50 basis points of expansion as software mix and productivity are offset by dilution from the early closing of Confluent. Our second quarter operating tax rate should be in the mid-teens. AI is fundamentally reshaping our clients' operating environments, increasing complexity, risk and the need for flexibility.

IBM's flywheel for growth built on trust, security and governance, a portfolio that helps enterprise put AI to work on their terms and sustained productivity that fuels rapid innovation, positions us to deliver value for our clients. We feel confident in our outlook and are excited about what's ahead. Arvind and I are now happy to take your questions. Olympia, let's get started.

Olympia McNerney: Thank you, Jim. Before we begin Q&A, I'd like to mention a couple of items. First, supplemental information is provided at the end of the presentation. And then second, as always, I'd ask you to refrain from multipart questions. Operator, let's please open it up for questions.

Operator: *[Operator Instructions]* And our first question comes from Amit Daryanani with Evercore ISI.

Amit Daryanani: Arvind, it's really nice to see the pickup in Red Hat growth and software acceleration broadly, especially as investors are debating software durability right now. When you step back and look at IBM's software portfolio, I want to understand how would you characterize your mix between infrastructure versus applications, consumption versus subscription kind of stuff. And as AI adoption really scales, where in that stack, do you see the most incremental value accruing to IBM versus the ecosystem? If you just frame how you think of software, the puts and takes in the AI-centric world, would be really helpful.

Arvind Krishna: Good to hear from you, Amit. And as you can imagine, that's the question that occupies us, actually occupies our clients, and I know it occupies many investors' minds. So let me first directly frame the answer in the dimensions that you laid out. If I think of infrastructure versus applications, I think if I count right, 4% of our portfolio, if I'm to be generous, could be called an application.

Specifically, I think the only part of our portfolio that is applications would be [Maximo] and even that, which is looking at maintenance and asset management operations, given that many people would not really call an application because it is the system of record as utilities and other people with very expensive infrastructure, keep their maintenance records, including where a part may have come from 30 years ago, 30 years ago. So why do I say that? So if you look at the Red Hat portfolio, that's operating systems and container-based software and automation and runbook software. I think we would correctly call that all enabling software.

The word 30 years ago would have been middleware, but that word has sort of gone away. If I look at our data portfolio, it is data basis, both relational and nonrelational, and then there is data movement and then there is AI enabling. But data movement is Confluent. And by AI, it's the Watson portfolio. In automation, it's all about helping people take complexity out of how they manage their IT infrastructure be that Turbonomic or Apptio or HashiCorp. And then there is mainframe software, which is largely very similar to the first 3 I mentioned, but for mainframe.

If you look at these, that is all I think in your words, I'll call it infrastructure, but I would call it more enabling software. Second part of your question was on subscription versus consumption. I think our entire portfolio is very tied to consumption. We sometimes use the word capacity because our mainframe, it is capacity, but capacity is equal to consumption. It's literally the MIPs that people use. It's not actually the installed capacity of the machine and off the mainframe in the distributed world, whether on the cloud or on-premise. A lot of it is sometimes tied to amount of compute capacity that the software runs on. That's consumption in a lot of sense.

Nobody is going to put it on processors if you're not actually consuming it. And that's the vast majority of our software portfolio. So it is very much tied to that. But then I think the implicit question you're asking is, why would that get a tailwind from AI as opposed to a headwind? So as people get serious, about AI because when they start experimenting, they may take a little bit of the data, they make a copy of it, they put on a public cloud, they run it on some public frontier model, they get some results, and that's exciting to them. As they get to scale, they've got to use the data from their internal systems.

If they're using data from the internal systems, many parts of our portfolio, be it Red Hat, be it Confluent, will come to be consumed more and more. As that gets consumed more and more, the automation part of the portfolio gets consumed more and more. As people do more and more fraud protection, not on sampling 1 in 10 transactions in the mainframe, but every single, that causes the mainframe consumption to go up. And we can see that, by the way, in the mainframe numbers we printed in the first quarter. So as we go through all of this, I think that this is a tailwind because of the model that we picked.

And by the way, I'll point out, this is not a model that's an accident of history. We have very consciously over the last 7 years driven the portfolio into this because we remain convinced that there is value in the underlying data layers. There is value in the business logic and then there's the interaction layer. Value is going to decrease in that interaction layer because as agents replace people for some fraction, we can debate how much of the interactions, then the interaction layer by itself is not sticky. The agents are going to be interacting much more with the underlying data and the business logic.

And we sort of saw that coming 6, 7 years ago, and that is why we picked the portfolio we did. I think that, that hopefully gives you the sense and we can see it in our numbers that AI is structurally increasing the demand for the portfolio, but that is also why both the organic products that we are building, for example, in software development and some of the acquisition targets we have had to play into the tailwinds of AI demand.

Operator: Our next question comes from Wamsi Mohan with Bank of America.

Wamsi Mohan: You said greater than 10% growth in software now in 2026. The early close of Confluent itself should add about a point of growth just by itself. So how are you seeing the growth trajectory for the remainder of the software portfolio as we go through 2026. And Arvind, maybe quickly for you, is IBM's appetite for M&A changing now that Confluent closes behind you and given the broader [derate] in the software space.

James Kavanaugh: Thanks, Wamsi, for the question. I appreciate it. Let's break down our software portfolio overall. First of all, we exit first quarter feeling very confident. In our software portfolio, the innovation value and the value proposition. And I think that goes to the core of Amit's question and how Arvind answered it about how we're playing central to the thesis of where enterprise software and AI come together that has always been predicated on our innovation strategy, our capital investment strategy and our M&A strategy. And I think what you see in the first quarter is a reflection of the diversification of our portfolio and the durability of our software model coming out of the first quarter.

But if you down back 90 days ago, what did Arvind and I say entering 2026, we felt very confident about the strategic repositioning of software. Why? One, portfolio shift to higher growth end markets; two, strong annuity base that we've been building up both organically and inorganically. By the way, we exit first quarter approaching $25 billion, growing 10%, [indiscernible] new innovation and GenAI realization, and we could talk about that, M&A growth synergies. And now we're into the second year of a very encouraging TP monetization opportunity. But for the full year, how is it going to play out? We talked about entering the year 10% growth. Now we see it growing 10-plus percent accelerating.

Data we started out extremely strong, growing 16%. We are taking data up for the year. Yes, we closed Confluent early, let's call it a couple of months overall, close it at the end of March. We were assuming somewhere in the mid-May time period in the end of second quarter. We now see data up low 20-plus percent range. That's going to deliver 5 points of software growth. That's representative of new innovation GenAI, the value of our platform-centric model and strategic partnerships and then also M&A contribution from Confluent, which should be about a little bit north of 15 points of that 20% to 25% growth overall for the year. So it's very strong organic. Hybrid cloud.

Red Hat entered the year, we accelerated, delivered as expected, we posted 10% growth. Underneath that, contributed 2.5 points to IBM for the year, and we're well positioned for that. Our subscription business accelerated. We got revenue under contract double digits, Red Hat OpenShift accelerating $2 billion ARR virtualization, now north of [indiscernible] and consumption model returned back to expectation. We are monitoring RHEL. RHEL did decelerate. I think that's a function of the federal lack of signings in the closure of the government in the fourth quarter that played through, but also a very dislocated hardware supply chain market. Automation on model, delivering over 2 points of growth.

Hashi great first year, record signings, we generated over $200 million in new incremental ARR that should position 2026 well, new innovation, M&A growth synergies and [indiscernible], continued growth and we're off to a tremendous start, record start in our new z17. So we actually feel very good and more optimistic than where we were 90 days ago on software.

Arvind Krishna: Let me just address the M&A question, Wamsi, very quickly. Yes, the values that are out there right now are very attractive. That does not always mean that the sellers are willing to accept these values that may take a few months for them to acknowledge that this is a new baseline. So if that's the case, I'm acknowledging that these are very attractive values. Now we have been a very disciplined acquirer one, let us make sure that we fully integrate in and get all the benefits from Confluent. So that is going to take some months to get done.

As we get through that and as the markets are at these values, that does open up our appetite perhaps more than it would in a normal year, but it's going to take a few months before we can go acknowledge whether or not that's going to happen. And that's where I would sort of give you the bit of color. So second half, if things stay where they are and if the values are where they are, maybe we can do something in the second half as we build up our cash balances, and we are 100% sure that Confluent is off to a strong start.

Operator: Our next question comes from Ben Reitzes with Melius Research.

Benjamin Reitzes: Appreciate it. Arvind and Jim, I just want to talk about guidance. You guys rarely raised guidance after first quarter, I get it. I think there's just some concern out there as to -- are you seeing something in Europe that keeps you at bay right now? Are you seeing evidence of something slowing that keeps you from raising guidance? I mean there's so many good things that are going on with regard to infrastructure and the software that you went through. So just wanted to clarify that -- and then with -- also with regard to guidance, the free cash flow was better than expected in the quarter. Why not raise it?

Or do you just need to see more evidence .

Arvind Krishna: Great. Ben, let me start out with just describing a little bit of the macro and what we have seen as evidence and then why we are being a little bit prudent. And then Jim will address all your questions on the specific guidance. Let me start with the Middle East. We had the strongest growth we have seen in decades, not years, decades in the Middle East. . So that gives you a sense that we are not seeing. There is no signal. I would tell you that I would expect the second quarter will play out similar to the first quarter in the Middle East.

Our clients there be it the larger enterprises, be it government, they are clear. They need to use and leverage technology to improve their own business. In the first quarter, Europe was also strong. You can see that in the supplemental materials that we have provided. So there is nothing in the -- what has already transpired. There has been no slower RN deals have actually progressed at the rate and pace that we would want. If I look at pipeline and demand signals of the second quarter, we are not seeing any of this slow down.

The only macro comment that we make is if the straits stay closed for another few weeks, then we know that there could be energy impacts in Europe, but that is speculative. That is not what we are seeing. And I expect that actually some of that we'll be able to absorb and maintain our acceleration. It's only if it crosses a certain level. So just based on only 3 months of the year have gone by is why we're making the prudent comment? Jim?

James Kavanaugh: Yes, Ben, thanks for the question. Let's take a step back and put this in perspective because I think you teed it up extremely well. I've been in this role now 9 years, Arvind's been in the role 6, 7 years. I don't think we've ever raised guidance in the first quarter. But let's talk about the mentality. We've done a lot of work about strategically repositioning our portfolio, our business operating model and the structural competitiveness of this business. And part of that was around how we were going to build discipline around execution in this company. And that execution mentality was around always [a beat] mentality at the end of the day.

The numbers speak for themselves in the first quarter. Strongest first quarter revenue growth that we've had in over a decade. Arvind talked about the macro environment. Arguably, yes, we're operating in a dynamic world. And there's more uncertainty than there was 90 days ago, as we all know. But within our lens of what we're looking at, we're executing extremely well across our high-value innovation software, infrastructure and consulting that see signs of progress. Underneath that, look at what's happening to the fundamentals of our business, our operating margins are up 140 basis points, our earnings are up nearly 20%, profits up 23%. This is an extremely strong start to the year.

And now you get the free cash flow, which I know is a valuation measure as I spent a lot of time out with our investors talking about our strategic narrative and our financial investment thesis, yes, Free cash flow generation is the multiple that people more and more are valuing IBM, and I would agree with that completely. We started out with the strongest free cash flow position in over a decade, highest free cash flow margin, up mid-teens. Let's put this in perspective. Less than 15% of what's required for the year. We got a lot of work ahead of us. But let's also put it in perspective, dial back a year ago, same call, same question.

Look at how we execute on that mentality that Arvind's been trying to drive in this company. We had that same discussion. We executed well. We took up free cash flow throughout the year, and then we blew through it in the fourth quarter. And I will tell you coming out of the first quarter, there's no different mentality that we have here today. The underlying fundamentals are adjusted EBITDA, by the way, all of this is high quality, sustainable, high-value realization overall. That is our free cash flow engine flywheel that provides tremendous investment flexibility for us to continue to invest and drive long-term sustainable competitive advantage, and we don't see any difference coming out of first quarter.

But again, first quarter in, we're 90 days into an extremely important year. And our view is we should be prudent.

Operator: Your next question comes from Fatima Boolani with Citigroup.

Fatima Boolani: Arvind, I wanted to pull on a thread in your prepared remarks with respect to the mainframe potentially being a destination for more emerging use cases, especially around AI inferencing. So call them not your traditional or conventional mainframe use cases, I was hoping you could put some quantitative framing around that. What type of a workload mix are you seeing today that you would consider conventionally mainframe? And what is that velocity of potential mix shift? And then as a related matter, as we think about the transaction processing and the MIPS growth momentum, how should that transpire and be expressed in the business in terms of the growth cadence for that particular segment of the follow-through.

I appreciate there's a little bit of a lag there. But would love your [indiscernible] and Jim's comments on that.

Arvind Krishna: Great. Thanks for the question, Fatima. Let me address the first part of your question, and then I'll actually give it to Jim to address some of the quantification of those workloads. So if we step back and look at it over the last 60 years mainframe has driven 2 great ways to monetize it. One has been what we call the classic MIPS or these are the compute parameters underneath that drive the transactional workloads that are great for the mainframe. Many people actually don't realize, but there are also, we call them Linux MIPs that are associated with the mainframe that people have been using to great effectiveness.

But let me acknowledge it is more sparse Linux workloads as opposed to the highly, highly intense Linux workloads. AI is adding a third kind of compute capacity into the mainframe. So just to make it very simple. Today, if people are doing a payment authorization, almost all the credit card companies in the world use the mainframe for their credit card authorizations. If they want to do fraud, they can run a few rules in that engine, but then they'll take a sampling of the transactions, let's call it, 10% of the platform because the latency that it introduces to take it off platform, you can't take them all, just slow the whole system down.

That's what they do off. What happens if you could run a 20 billion, 30 billion parameter model right on the mainframe, suddenly because that is only milliseconds of latency, you can do that to every single transaction. So if you can take your fraud rate down from 50 basis points to 40, you can now do the math on what that is. They are all seeing that. So as I do that, I think we can do that for credit card authorizations. We can do it for retail banking transactions. We can do it for other payment operations. We can do it for claims and billing purposes. So those are the workloads that are now coming on.

So it is effectively a new capacity of the mainframe that previously was either very small but outside the mainframe or running on systems that are what we would call distributed infrastructure. We believe that this is going to play out. We see a large majority of our clients asking for the capacity. And currently, I believe we have a fully populated system we can do about 450 billion inferences a day on the mainframe. So that gives you a sense of that. We monetize that both through the extra hardware that is sold but also by the supporting software for all of the AI inferencing that then runs on that increased capacity.

So with that then, hopefully, that gives you some color on what is happening. I'll give it to Jim.

James Kavanaugh: Yes. Thank you for the question overall. I mean, mainframe modernization increases the strategic importance of IBM z, as Arvind talked about. Why? Because the source of value is architectural. It's the platform. It's not the language. It's a tight integration of software, hardware, database, security, run time, resiliency. And as Arvind talked about, this is a whole new monetization area of opportunity for us on that platform stack. What is the driver of growth?

Yes, 450 billion AI inferences at 1 millisecond of response time, 25 billion encryptions, transactions per day, up to eight 9s of availability, quantum-safe encryption and a TCO advantage running it on mainframe, on-prem versus the cloud anywhere from 3 to 15x depending on the size and complexity of that platform. So that's why mainframe runs 73% of the world's transaction volumes in terms of value, 45 of the top 50 banks, 9 of the top 10 retailers, 4 to 5 top airlines, et cetera. Now you go to your second question about how do we monetize that value. One is the monetization of the platform of hardware Arvind talked about AI MIPS. Second is that stack economic multiplier.

Historically, we've been averaging about 3x to 4x stack multiplier for every hardware dollar we land. Let me give you a stat. We just anniversaried our first full year of z17. That first full year is z17 versus the prior program, z16 first full year, which, by the way, was the best on record at that point in time. We've increased hardware placement value by over $1 billion. Now you take that $1 billion and you think about the future monetization opportunity that we get. That's that 3 to 4x multiplier that will play out over time. A big chunk of that being our TP software, but it's also our storage attach, it's our maintenance business, it's our financing business.

We monetize that value based on how many MIPS are shipped in the market and for 4 quarters in a row, on Z17, we've shipped over 100% growth of new MIPS in the market, including first quarter. Why does that matter? Higher capacity is higher monetization opportunity, it's higher price opportunity, it's higher value creation opportunity. So we feel pretty good about that future modernization and multiplier effect as we play out 2026 and 2027.

Operator: Our next question comes from Brent Thill with Jefferies.

Brent Thill: Jim, just on the constant currency for software, not to nitpick, but if you look at last year, 9% growth in Q1, I think it was 8% -- 11% in Q4. So investors are asking, you're seeing a little bit of a downtick. Is that due to seasonality where maybe your contract signings were better in Q1, but maybe are being reflected in the reported numbers? Again, I know it's a modest deceleration, but anything to point out there?

James Kavanaugh: Yes, Brent, thank you very much for the question overall. I fully expected this one because when you just look at the media, print and the press release, fourth quarter, we posted a little over 11% growth. This quarter, we posted 8% growth. What gives -- do you feel still strength about your portfolio, your business, your investments, your new innovation, I think you nailed it right upfront. One, understanding our business, our software portfolio high-value recurring revenue, about 80% of our annual business, about $30 billion plus trailing 12 months, 80% of that is high-value annuity-based business, is a transactional engine underneath it.

It's a big component of our perpetual license model, but it's a component of our subscription model, et cetera. If you look at it. The entire 3-point drop quarter-to-quarter is the fundamentals of the mix of the portfolio. In fourth quarter, we have about 30% of our business in the fourth quarter is transactional. In the first quarter, that's about 10%. When you look at the underpinnings of the core annuity by itself, we're actually accelerating that fourth quarter to first quarter. I think I said earlier on the call, our annuity ARR exiting first quarter approaching $25 billion, that's up 10%.

Throughout the rest of the year, we'll go from a transactional quarter of about 10% first and will peak probably in the fourth quarter of about 30-plus percent. On average, we'll be in the 20% overall. That will accelerate growth. That, coupled with M&A growth synergies, our GenAI portfolio, which has had a lot of momentum behind it. And our TP monetization and cycle I would tell you, coming out of first quarter, I feel pretty good about 8% growth, and it positions us why we said 90 days ago, confident in 10%. Now we're saying, yes, we closed Confluent earlier, and we're confident now in accelerating that to 10 plus.

Operator: Our next question comes from Erik Woodring with Morgan Stanley.

Erik Woodring: Jim, you briefly alluded to it earlier, but -- can you maybe just detail how IBM is broadly managing and/or mitigating some of these supply chain headwinds, whether that's higher memory costs or supply challenges, meaning how material is memory within the infrastructure base? How are you mitigating. How are customers responding how does it impact your outlook on growth and margins? If you could just maybe dig into this, that would be super helpful. .

James Kavanaugh: Yes, absolutely. Thank you, Erik, for the question overall. You understand our business extremely well. Underneath Arvind's leadership, we have strategically reposition this portfolio. There's been a lot of work around portfolio optimization. By the way, that's both leveraging the strength of our cash flow, our financial flexibility, to buy high-value innovative base companies in category-leading technologies with structural growth profiles to help IBM but it's also around divestitures of portfolio. But where I'm going with this is today, when you look at IBM's portfolio, we're a human capital asset, IP-based business at 75% on a way to 80% by the way, underneath that, software, 45% on its way to 50-plus percent. Overall.

Our hardware business is extremely important as a value creator to IBM -- but top line, it's about 25% of our business, but that's high-value innovation on mainframe platform overall. Now you look underneath it around the supply chain dislocation around, commodity cost increases, in particular around memory, it has a de minimis impact to us overall. Think about mainframe overall. Will it impact storage and potentially some components of our distributed infrastructure, absolutely. But look underneath it, we're able to -- one, we've been in existence for 115, 116 years overall. We know how to run global supply chains. We drive supplier optimization, supply chain diversification, procurement strategies overall.

And I think we've been able to mitigate this dislocation overall. The area we're watching it is in the software area around RHEL. I mean RHEL's tied to enterprise hardware placements overall, and we'll continue monitoring that. But look at our hardware performance, we accelerated growth at 15%, our distributed infrastructure at actual rates growing 17%, constant currency growing 13%. By the way, I didn't even talk about it. our infrastructure pretax margins are up 720 basis points. So we know how to manage global supply chains and commodity costs inside the company and extract value overall.

Arvind Krishna: Jim, let me just add a couple of sentences to your statement. Erik, Jim mentioned that we have worked with a lot of the suppliers for a number of years and decades. They like working with us, partly because of the relationships we have built up with them over the years, but also because we help them stress test new capabilities and they like the fact that our systems are very high performing because that gives them brand reputation as they go out of the wider market. That does not help completely, but somewhat mitigate some of the supply chain constraints because we are early users of the [new S memory] technologies.

Operator: Your next question comes from Jim Schneider with Goldman Sachs.

James Schneider: I was wondering if you could maybe comment on the AI bookings, which is a metric you previously given -- you've given, but I think you just commented as a percentage of your total bookings right now. Did that accelerate or decelerate in the quarter? And then maybe just kind of comment on any update you see for the consulting business this year, either given more macro uncertainty. Do you expect any kind of diminution in the growth rate you expect this year? .

James Kavanaugh: Yes. Thanks, Jim. As we talked about, we exited last year with a book of business around AI, which, as you know, we talked about consulting and software within that vernacular. I think it was important over the last couple of years because as the explosion of GenAI hit, we had to give a perspective about whether we were winning and capitalizing and participating in that market. We exited last year what $12.5 billion, over $12.5 billion book of business. But now let's bring it back because in January, we talked about it's embedded across our portfolio. It's embedded in software, Its central thesis to how we run our consulting business right now. It's embedded across our infrastructure business.

And we said coming into 2026, we were going to talk about it more from an outcome-based revenue base and value contribution base overall. But let's talk about software. Software GenAI continues to be a tailwind overall. The positioning of our portfolio with the explosion of AI, with applications agents with us owning the foundational layer of Linux, containerization, you see that play out with the acceleration Red Hat OpenShift business, now $2 billion, growing north, I think, high 20% growth overall. Second, the importance of the data layer. Arvind talked about Confluent positioning us to be the cross platform as a data connector, automation, the need of resiliency, observability FinOps.

Software, let's talk about, one, it's accelerating our growth profile overall. But let me put some numbers behind it versus just an overarching book of business. Our software book from an annualized revenue trailing 12 months, we finished last year at $30 billion, right? 80% of that, as I said earlier, high-value recurring revenue, 20% transactional. We did about $6 billion. Over the last trailing 12 months on an accelerating basis, our AI platform agents, assistance orchestration is north of $1.5 billion. It's already about 25% penetrated and our software business growing north of 40%. It's contributing 2 points of growth on an annualized basis. And a thing we love about it, it has a multiplier effect over time.

So it's an acceleration there. Consulting. Consulting is about 40% of our signings, 30% of our backlog is GenAI now, over 20% of our revenue. And on an ARR revenue perspective, in the first quarter, we eclipsed $4 billion ARR. So it is central to the way we run a services as software model overall. And in infrastructure, both Arvind and I talked about, it's embedded on the chip of z17, [indiscernible] inferencing. I think Arvind talked about it in prepared remarks, clients that have implemented watson Code assistant for Z, we're seeing 3x differential on growth and capacity, and you see in our distributed infrastructure, we're accelerating growth. Now your second question around consulting.

We are seeing signposts of progress overall. One, our demand profile, our backlog quality, our GenAI, which I just talked about, our strategic partnership headroom opportunity, our portfolio mix composition more to higher growth areas and our services as software model, which we think we have an industry-leading position with our IBM Consulting Advantage platform. But let me put some stats on it. One, signings, we return to growth. Great quarter overall, large transformational deals around GenAI, the health and mix of net new business and expansions up 7 points year-to-year, up 4 points quarter-to-quarter. 400 new clients captured in the first quarter. Our backlog quality overall, our erosion is stable. Our duration continues to come down.

Our backlog realization is actually accelerating throughout the year, and our backlog yields are up 4 points year-over-year talking about the quality and value we're able to deliver. And I talked about GenAI, 80% of our GenAI book of business right now is coming from capture from net new clients overall. And I'll stress that over $4 billion revenue ARR. So that positions our confidence in the year of us accelerating our revenue growth around low single digits and if things go well, can we do better than that? Obviously, yes.

Olympia McNerney: Operator, let's take 1 last question.

Operator: Last question comes from Matt Swanson with RBC Capital Markets.

Matthew Swanson: Great. Thank you so much for squeezing me in here. Arvind, it's really interesting going over the software segments, and you showed how [low] of an exposure you have to the application space. There's obviously been a ton of debate right now around who's going to kind of win the GenAI workloads of application. We've seen you operate at such a strong kind of [Switzerland] foundational player in the hybrid cloud. When we look at AI, like how are you setting IBM up to win kind of regardless what ends up being the winner of the GenAI application layer? And I mean, what kind of investments does that take?

Arvind Krishna: Matt, thanks for the question. So we made the decision about 3 years ago that we were going to be neutral and Switzerland like also on our usage of frontier models. Because I think when we are saying the GenAI applications, I think for many people that is synonymous with the frontier model providers, not just the fronter models, but all the surrounding software [indiscernible] that all of them are giving. So we are going to play where clients want to be hybrid.

The clients want to function across multiple clouds or also because of either sovereignty or brand or privacy or in the end, economics, they might also have a private addition in addition to what they use on public. So as we go across that, we are building, for example, our software development AI product, Project Bob. It is out, we actually chose not to announce it. Nevertheless, 200 people signed up to use it. So that gives us a signal that we have something.

Now why would they use us as opposed to just one of the Codex or equivalent models is if they also have a lot of code that they do not want to actually take out in public. And also, they want to address the entire software development, meaning including testing, including [patching], including documentation, including maintenance are the kinds of things that we provide. Ditto as we look at how they might want to use agents that come inside their enterprise, then we use Confluent to go manage and control how they expose data from inside things. So as we sort of look at that math, I think we're very clear right.

There will be people who will be frontier model providers. You can debate are those half a dozen or a dozen? Today, it's somewhere in that range. We actually do not want to even predict which of them will be the eventual winners. We want to work with all of them. Then we also work with open weight models. And we produce models where we have either domain expertise or people may want much smaller models to be able to run them on-premise or I'll say, euphemistically on a 1 to 4 GPU server node as opposed to a very, very large cluster. So that's the model place.

We are going to then help our clients deploy these models to gain value. As we have unlocked, Jim talked about the $4.5 billion of internal value, how do you reduce your total tax expense? How do you reduce procurement expense? How do you reduce accounts payable, how do you reduce [quote to cash] as we walk across these processes, we get a lot of knowledge on how to capture that into agents, but then we are not going to be fixated whichever model you want to use, you can use. And wherever you want to run them, we'll help you run them.

And we think that's a good half of the world is interested in that paradigm, and that's how, Matt, we are going to be able to go win in this world as it unfolds going forward. So just to close, the innovation value we are delivering to our clients and our strong start to the year, reinforce our confidence in our growth trajectory. We look forward to continuing this dialogue as we move through the year.

Olympia McNerney: Thank you, Arvind. Operator, let me turn it back to you to close out the call.

Operator: Thank you for participating on today's call. The conference has now ended. You may disconnect at this time.

Read Next

Jul 13, 2026 By Rick Orford
Great News for IBM Stock Investors!

Jul 9, 2026 By Scott Levine
7 Best Cloud Computing Stocks for 2026 and How to Invest

Jul 8, 2026 By Anders Bylund
9 Best Quantum Computing Stocks for 2026 and How to Invest

Jul 2, 2026 By Rick Orford
This Tech Stock Could Benefit From the Next Computing Breakthrough

Jul 1, 2026 By Rick Orford
Quantum Stocks Face a 2028 Cash Test. Here's What Investors Should Know.

Jun 25, 2026 By Micah Zimmerman
The U.S. Government Is Betting Billions on Quantum Computing. These 3 Stocks Are the Biggest Winners.

This article is a transcript of this conference call produced for The Motley Fool. While we strive for our Foolish Best, there may be errors, omissions, or inaccuracies in this transcript. Parts of this article were created using Large Language Models (LLMs) based on The Motley Fool's insights and investing approach. It has been reviewed by our AI quality control systems. Since LLMs cannot (currently) own stocks, it has no positions in any of the stocks mentioned. As with all our articles, The Motley Fool does not assume any responsibility for your use of this content, and we strongly encourage you to do your own research, including listening to the call yourself and reading the company's SEC filings. Please see our Terms and Conditions for additional details, including our Obligatory Capitalized Disclaimers of Liability.
The Motley Fool has positions in and recommends International Business Machines. The Motley Fool has a disclosure policy.

Premium Investing Services

Invest better with The Motley Fool. Get stock recommendations, portfolio guidance, and more from The Motley Fool's premium services.

View Premium Services

Making the world smarter, happier, and richer.

© 1995 - 2026 The Motley Fool. All rights reserved.
Market data powered by Xignite and Polygon.io.

ABOUT THE MOTLEY FOOL
About Us
Careers
Research
Newsroom
Contact
Advertise

OUR SERVICES
All Services
Stock Advisor
Epic
Epic Plus
Fool Portfolios
Fool One
Motley Fool Money

AROUND THE GLOBE
Fool Australia
Fool Canada

FREE TOOLS
CAPS Stock Ratings
Discussion Boards
Calculators
Financial Dictionary

AFFILIATES & FRIENDS
Motley Fool Asset Management
Motley Fool Wealth Management
Motley Fool Ventures
Fool Community Foundation
Become an Affiliate Partner

Terms of Use | Privacy Policy | Disclosure Policy | Accessibility Policy | Copyright, Trademark and Patent Information | Terms and Conditions | Do Not Sell My Personal Information